

EE5179 : Deep Learning for Imaging

Programming Assignment 3: Autoencoders

Name: SRIVATSAN SARVESAN

Roll No: DA24E001

```
In [2]: import numpy as np
import torch
import matplotlib.pyplot as plt
import torchvision
import torchvision.transforms as transforms
from torch import nn
```

```
In [3]: #HYPERPARAMETERS
learning_rate=0.001
epochs=10
batch_size= 256
number_of_pc=30 #number of principal components (for PCA)
```

```
In [4]: train_data = torchvision.datasets.MNIST(root=".",train=True,transform=transforms.T
test_data = torchvision.datasets.MNIST(root=".",train=False,transform=transforms.T
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>

Failed to download (trying next):

<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: certificate has expired (_ssl.c:1007)>

Downloading <https://ossci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte.gz>

Downloading <https://ossci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte.gz> to ./MNIST/raw/train-images-idx3-ubyte.gz

100%|██████████| 9.91M/9.91M [00:00<00:00, 11.2MB/s]

Extracting ./MNIST/raw/train-images-idx3-ubyte.gz to ./MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz>

Failed to download (trying next):

<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: certificate has expired (_ssl.c:1007)>

Downloading <https://ossci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte.gz>

Downloading <https://ossci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte.gz> to ./MNIST/raw/train-labels-idx1-ubyte.gz

100%|██████████| 28.9k/28.9k [00:00<00:00, 351kB/s]

Extracting ./MNIST/raw/train-labels-idx1-ubyte.gz to ./MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz

Failed to download (trying next):

<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: certificate has expired (_ssl.c:1007)>

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz to ./MNIST/raw/t10k-images-idx3-ubyte.gz

100%|██████████| 1.65M/1.65M [00:00<00:00, 2.77MB/s]

Extracting ./MNIST/raw/t10k-images-idx3-ubyte.gz to ./MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz

Failed to download (trying next):

<urlopen error [SSL: CERTIFICATE_VERIFY_FAILED] certificate verify failed: certificate has expired (_ssl.c:1007)>

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz to ./MNIST/raw/t10k-labels-idx1-ubyte.gz

100%|██████████| 4.54k/4.54k [00:00<00:00, 8.18MB/s]

Extracting ./MNIST/raw/t10k-labels-idx1-ubyte.gz to ./MNIST/raw

```
In [5]: train_loader = torch.utils.data.DataLoader(dataset=train_data, batch_size=batch_size,
test_loader = torch.utils.data.DataLoader(dataset=test_data, batch_size=len(test_data))
```

```
In [6]: test_sample_loader = torch.utils.data.DataLoader(dataset=test_data.data[9705:9715],
```

```
In [7]: train_dataset=train_data.data.reshape(train_data.data.shape[0], train_data.data.shape[1], train_data.data.shape[2])
test_dataset = test_data.data.reshape(test_data.data.shape[0], test_data.data.shape[1], test_data.data.shape[2])
test_dataset_sampled= test_dataset[np.arange(9705, 9715), :]
```

This is formatted as code

Comparing PCA and Autoencoders

```
In [ ]: def PCA(input_data,top_k_ev):
    """
    function: This function performs pca.
    Input: input_data = (torch matrix) = of shape num_datapts,784pixels
           top_k_ev = (integer) = how many principal components to be taken

    Output: top_k_eigen_vectors= shape(784,top_k_ev) top k eigen vectors in columns
           centered_ip_data= shape(num_datapts,784) = centered ip data (used for rec

    """

    input_mean = torch.mean(input_data,0)
    centered_ip_data = input_data-input_mean
    cov_matrix = torch.matmul(centered_ip_data.T,centered_ip_data)
    eigen_values, eigen_vectors = torch.linalg.eigh(cov_matrix)
    eigen_values_descending,indices = torch.sort(eigen_values,descending=True)
    top_k_eigen_values,top_k_indices = eigen_values_descending[:top_k_ev],indices[:top_k_ev]
    top_k_eigen_vectors = eigen_vectors[:,top_k_indices]

    assert top_k_eigen_vectors.shape == (784,top_k_ev)
    assert centered_ip_data.shape == input_data.shape
    assert input_mean.shape == torch.Size([784])

    return top_k_eigen_vectors
```

```
In [ ]: pc=PCA(train_dataset.float(),number_of_pc)
```

```
In [ ]: def reconstruct_data(principal_components,dataset):
    """
    function: This function reconstructs the datapoints in lower dimension (i.e. top

    Input: principal_components= shape(784,top_k_ev) top k eigen vectors in columns
           centered_ip_data= shape(num_datapts,784) = centered ip data (used for rec

    Output: projected_data = (torch matrix) = of shape num_datapts , top_k_ev

    """

    projection_matrix = torch.matmul(principal_components,principal_components.T)
    projected_data = torch.matmul(dataset,projection_matrix)

    assert projected_data.shape == (dataset.shape[0],principal_components.shape[0])

    return projected_data
```

```
In [ ]: reconstructed_test_data_sampled=reconstruct_data(pc,test_dataset_sampled.float())
```

```
In [ ]: class AE1(nn.Module):
    def __init__(self):
        super(AE1, self).__init__()
        self.encoder = nn.Sequential(
            nn.Linear(784,512),
            nn.ReLU(),
            nn.Linear(512,256),
            nn.ReLU(),
            nn.Linear(256,128),
            nn.ReLU(),
            nn.Linear(128,30),
            nn.ReLU())
        self.decoder = nn.Sequential(
            nn.Linear(30,128),
            nn.ReLU(),
            nn.Linear(128,256),
            nn.ReLU(),
            nn.Linear(256,784),
            nn.ReLU())

    def forward(self,x):
        x=self.encoder(x)
        encoded_output=x
        x=self.decoder(x)
        return x,encoded_output
```

```
In [ ]: model1 = AE1()
criterion1 = nn.MSELoss()
optimizer1 = torch.optim.Adam(model1.parameters(), lr=learning_rate)

training_loss = []

epochs = 10
for epoch in range(epochs):
    epoch_loss = 0 # To accumulate loss over the entire epoch
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1) # Flatten the images
        outputs, _ = model1(images)
        loss = criterion1(outputs, images) # Calculate the loss

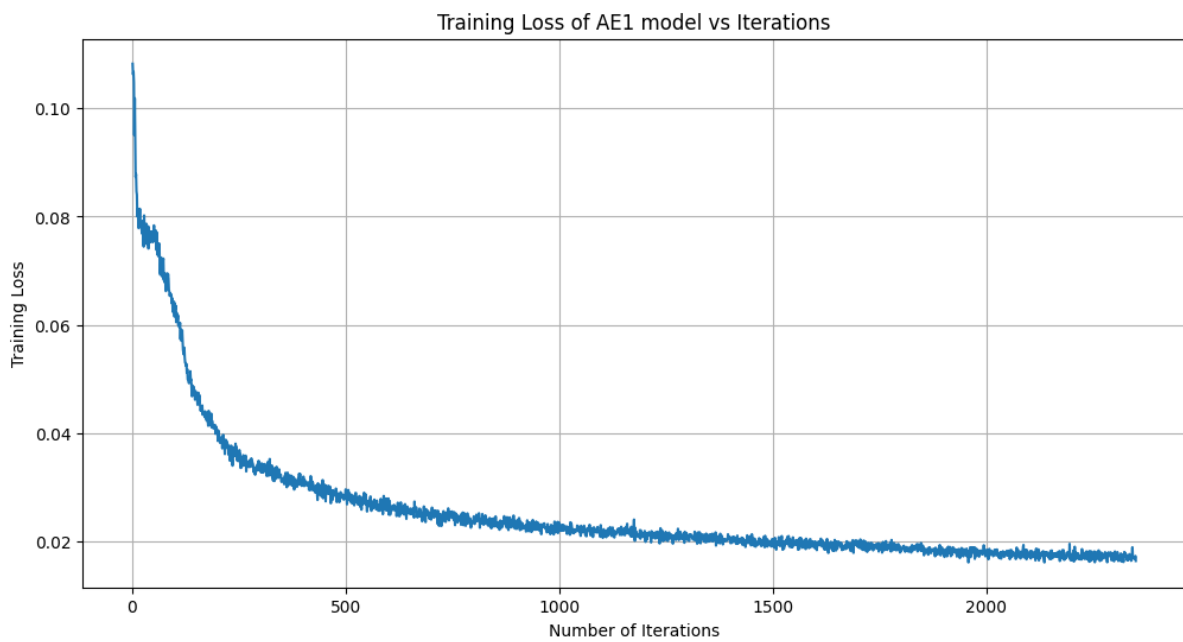
        training_loss.append(loss.item()) # Record the loss
        epoch_loss += loss.item() # Accumulate loss for the epoch

        optimizer1.zero_grad() # Zero the gradients
        loss.backward() # Backpropagation
        optimizer1.step() # Update the weights

    avg_loss = epoch_loss / len(train_loader) # Average loss for the epoch
    print(f"Epoch [{epoch + 1}/{epochs}], Loss: {avg_loss:.4f}") # Print the loss
```

```
Epoch [1/10], Loss: 0.0584  
Epoch [2/10], Loss: 0.0321  
Epoch [3/10], Loss: 0.0266  
Epoch [4/10], Loss: 0.0237  
Epoch [5/10], Loss: 0.0220  
Epoch [6/10], Loss: 0.0208  
Epoch [7/10], Loss: 0.0197  
Epoch [8/10], Loss: 0.0189  
Epoch [9/10], Loss: 0.0178  
Epoch [10/10], Loss: 0.0173
```

```
In [ ]: plt.figure(figsize=(12,6))  
plt.plot(range(1,len(training_loss)+1),training_loss)  
plt.title("Training Loss of AE1 model vs Iterations")  
plt.xlabel("Number of Iterations")  
plt.ylabel("Training Loss")  
plt.grid()
```



```

In [ ]: #MODEL EVALUATION AND RESULT PLOTTING
model1.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        outputs,_ = model1(images.float())

plt.rcParams["figure.figsize"] = (12,6)
for i in range (10):
    fig, (ax1, ax2, ax3) = plt.subplots(1,3)

    ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs[i].detach().numpy().reshape(28,28),cmap='gray')
    ax2.set_title('AE Reconstructed Image')
    ax2.axis("off")


    ax3.imshow(reconstructed_test_data_sampled[i].reshape(28,28),cmap='gray')
    ax3.set_title('PCA Reconstructed Image')
    ax3.axis("off")
    print("_____")
    print("Reconstruction Error in AE:",np.dot(((images[i].detach().numpy()/255)-(out
    print("Reconstruction Error in PCA:",np.dot(((images[i].detach().numpy()/255)-(re
    plt.show()

```


Reconstruction Error in AE: 6.307659293984471

Reconstruction Error in PCA: 4.906910216143052


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 21.927312816323273

Reconstruction Error in PCA: 16.862376030476067


Original Image



AE Reconstructed Image



PCA Reconstructed Image




Reconstruction Error in AE: 20.272061440857204

Reconstruction Error in PCA: 16.0457551172948


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 13.899205078610088

Reconstruction Error in PCA: 10.714796514097081


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 12.313258415085311

Reconstruction Error in PCA: 14.848193155147694


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 25.932653129596723

Reconstruction Error in PCA: 20.39081935165547


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 15.285770999594304

Reconstruction Error in PCA: 11.735863218712451


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 13.924664436835346

Reconstruction Error in PCA: 14.356571254531701


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 14.230142970600042

Reconstruction Error in PCA: 13.215490057306457


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 18.62658867292707

Reconstruction Error in PCA: 16.605797151615285


Original Image



AE Reconstructed Image



PCA Reconstructed Image



```
In [ ]: epochs=20

model2 = AE1()
criterion1 = nn.MSELoss()
optimizer1 = torch.optim.Adam(model2.parameters(), lr=learning_rate)

training_loss = []

for epoch in range(epochs):
    epoch_loss = 0 # To accumulate loss over the entire epoch
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1) # Flatten the images
        outputs, _ = model2(images)
        loss = criterion1(outputs, images) # Calculate the loss


        training_loss.append(loss.item()) # Record the loss
        epoch_loss += loss.item() # Accumulate loss for the epoch

        optimizer1.zero_grad() # Zero the gradients
        loss.backward() # Backpropagation
        optimizer1.step() # Update the weights

    avg_loss = epoch_loss / len(train_loader) # Average loss for the epoch
    print(f"Epoch [{epoch + 1}/{epochs}], Loss: {avg_loss:.4f}") # Print the loss
```

```
Epoch [1/20], Loss: 0.0623
Epoch [2/20], Loss: 0.0339
Epoch [3/20], Loss: 0.0277
Epoch [4/20], Loss: 0.0246
Epoch [5/20], Loss: 0.0229
Epoch [6/20], Loss: 0.0214
Epoch [7/20], Loss: 0.0203
Epoch [8/20], Loss: 0.0193
Epoch [9/20], Loss: 0.0187
Epoch [10/20], Loss: 0.0182
Epoch [11/20], Loss: 0.0177
Epoch [12/20], Loss: 0.0173
Epoch [13/20], Loss: 0.0169
Epoch [14/20], Loss: 0.0166
Epoch [15/20], Loss: 0.0161
Epoch [16/20], Loss: 0.0159
Epoch [17/20], Loss: 0.0157
Epoch [18/20], Loss: 0.0156
Epoch [19/20], Loss: 0.0154
Epoch [20/20], Loss: 0.0153
```

```
In [ ]: plt.figure(figsize=(12,6))
plt.plot(range(1,len(training_loss)+1),training_loss)
plt.title("Training Loss of AE1 model vs Iterations")
plt.xlabel("Number of Iterations")
plt.ylabel("Training Loss")
plt.grid()
```



```

In [ ]: #MODEL EVALUATION AND RESULT PLOTTING
model1.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        outputs_1,_ = model1(images.float())

model2.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        outputs_2,_ = model2(images.float())

plt.rcParams["figure.figsize"] = (20,6)
for i in range (10):
    fig, (ax1, ax2, ax3, ax4) = plt.subplots(1,4)

    ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")


    ax2.imshow(outputs_1[i].detach().numpy().reshape(28,28),cmap='gray')
    ax2.set_title('AE Reconstructed Image(epochs=10)')
    ax2.axis("off")

    ax3.imshow(outputs_2[i].detach().numpy().reshape(28,28),cmap='gray')
    ax3.set_title('AE Reconstructed Image(epochs=20)')
    ax3.axis("off")


    ax4.imshow(reconstructed_test_data_sampled[i].reshape(28,28),cmap='gray')
    ax4.set_title('PCA Reconstructed Image')
    ax4.axis("off")
    print("
    print("Reconstruction Error in AE(epochs=10):",np.dot(((images[i].detach().numpy()
    print("Reconstruction Error in AE(epochs=20):",np.dot(((images[i].detach().numpy()
    print("Reconstruction Error in PCA:",np.dot(((images[i].detach().numpy()/255)-(re
    plt.show()

```


Reconstruction Error in AE(epochs=10): 6.307659293984471
 Reconstruction Error in AE(epochs=20): 7.201358925435635
 Reconstruction Error in PCA: 4.906910216143052




Reconstruction Error in AE(epochs=10): 21.927312816323273
Reconstruction Error in AE(epochs=20): 30.92957060781849
Reconstruction Error in PCA: 16.862376030476067




Reconstruction Error in AE(epochs=10): 20.272061440857204
Reconstruction Error in AE(epochs=20): 15.765950224609877
Reconstruction Error in PCA: 16.0457551172948




Reconstruction Error in AE(epochs=10): 13.899205078610088
Reconstruction Error in AE(epochs=20): 12.61424778484006
Reconstruction Error in PCA: 10.714796514097081




Reconstruction Error in AE(epochs=10): 12.313258415085311
Reconstruction Error in AE(epochs=20): 12.80333184312766
Reconstruction Error in PCA: 14.848193155147694




Reconstruction Error in AE(epochs=10): 25.932653129596723
Reconstruction Error in AE(epochs=20): 27.042785428365406
Reconstruction Error in PCA: 20.39081935165547




Reconstruction Error in AE(epochs=10): 15.285770999594304
Reconstruction Error in AE(epochs=20): 13.609016259185033
Reconstruction Error in PCA: 11.735863218712451




Reconstruction Error in AE(epochs=10): 13.924664436835346
Reconstruction Error in AE(epochs=20): 15.198742691518415
Reconstruction Error in PCA: 14.356571254531701



Reconstruction Error in AE(epochs=10): 14.230142970600042
Reconstruction Error in AE(epochs=20): 14.036274052054509
Reconstruction Error in PCA: 13.215490057306457



Reconstruction Error in AE(epochs=10): 18.62658867292707
 Reconstruction Error in AE(epochs=20): 20.33213802100416
 Reconstruction Error in PCA: 16.605797151615285



Observations:

- Visually, Reconstructed Autoencoder images look appealing as they have better contrast.
- Reconstruction error-wise, Reconstructed PCA images are found to have lesser error (doesnot have much difference though), however they lack the contrast.

Standard Autoencoder

```
In [ ]: class AE2(nn.Module):
    def __init__(self,hid):
        super(AE2, self).__init__()
        self.hid=hid
        self.encoder = nn.Sequential(
            nn.Linear(784,self.hid),
            nn.ReLU()
        )
        self.decoder =nn.Sequential(
            nn.Linear(self.hid,784),
            nn.ReLU()
        )

    def forward(self,x):
        x=self.encoder(x)
        encoded_output=x
        x=self.decoder(x)
        return x,encoded_output
```

```
In [ ]: # Model with hidden size = 64
epochs = 10
model_hid64 = AE2(64)
criterion_hid64 = nn.MSELoss()
optimizer_hid64 = torch.optim.Adam(model_hid64.parameters(), lr=learning_rate)

training_loss_hid64 = []

for epoch in range(epochs):
    epoch_loss = 0 # To accumulate the loss over the epoch
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1) # Flatten the images
        outputs, _ = model_hid64(images)
        loss = criterion_hid64(outputs, images) # Calculate Loss

        training_loss_hid64.append(loss.item()) # Record loss for each batch
        epoch_loss += loss.item() # Accumulate the loss for the epoch

        optimizer_hid64.zero_grad() # Zero the gradients
        loss.backward() # Backpropagate the loss
        optimizer_hid64.step() # Update the weights

    avg_loss = epoch_loss / len(train_loader) # Compute average loss for the epoch
    print(f"Epoch [{epoch + 1}/{epochs}], Loss: {avg_loss:.4f}") # Print average L
```

```
Epoch [1/10], Loss: 0.0364
Epoch [2/10], Loss: 0.0164
Epoch [3/10], Loss: 0.0140
Epoch [4/10], Loss: 0.0130
Epoch [5/10], Loss: 0.0124
Epoch [6/10], Loss: 0.0120
Epoch [7/10], Loss: 0.0118
Epoch [8/10], Loss: 0.0117
Epoch [9/10], Loss: 0.0116
Epoch [10/10], Loss: 0.0115
```



```
In [ ]: # Model with hidden size = 128
model_hid128 = AE2(128)
epochs=10
criterion_hid128 = nn.MSELoss()
optimizer_hid128 = torch.optim.Adam(model_hid128.parameters(), lr=learning_rate)

training_loss_hid128 = []

for epoch in range(epochs):
    epoch_loss = 0 # To accumulate the loss over the epoch
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1) # Flatten the images
        outputs, _ = model_hid128(images)
        loss = criterion_hid128(outputs, images) # Calculate loss

        training_loss_hid128.append(loss.item()) # Record loss for each batch
        epoch_loss += loss.item() # Accumulate the loss for the epoch

        optimizer_hid128.zero_grad() # Zero the gradients
        loss.backward() # Backpropagate the loss
        optimizer_hid128.step() # Update the weights

    avg_loss = epoch_loss / len(train_loader) # Compute average loss for the epoch
    print(f"Epoch [{epoch + 1}/{epochs}], Loss: {avg_loss:.4f}") # Print average L
```

```
Epoch [1/10], Loss: 0.0270
Epoch [2/10], Loss: 0.0110
Epoch [3/10], Loss: 0.0092
Epoch [4/10], Loss: 0.0085
Epoch [5/10], Loss: 0.0081
Epoch [6/10], Loss: 0.0078
Epoch [7/10], Loss: 0.0077
Epoch [8/10], Loss: 0.0076
Epoch [9/10], Loss: 0.0075
Epoch [10/10], Loss: 0.0074
```

```
In [ ]: # Model with hidden size = 256
model_hid256 = AE2(256)
criterion_hid256 = nn.MSELoss()
optimizer_hid256 = torch.optim.Adam(model_hid256.parameters(), lr=learning_rate)
epochs=10
training_loss_hid256 = []

for epoch in range(epochs):
    epoch_loss = 0 # To accumulate the loss over the epoch
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1) # Flatten the images
        outputs, _ = model_hid256(images)
        loss = criterion_hid256(outputs, images) # Calculate Loss


        training_loss_hid256.append(loss.item()) # Record Loss for each batch
        epoch_loss += loss.item() # Accumulate the loss for the epoch

        optimizer_hid256.zero_grad() # Zero the gradients
        loss.backward() # Backpropagate the loss
        optimizer_hid256.step() # Update the weights

    avg_loss = epoch_loss / len(train_loader) # Compute average loss for the epoch
    print(f"Epoch [{epoch + 1}/{epochs}], Loss: {avg_loss:.4f}") # Print average L
```

```
Epoch [1/10], Loss: 0.0228
Epoch [2/10], Loss: 0.0099
Epoch [3/10], Loss: 0.0086
Epoch [4/10], Loss: 0.0076
Epoch [5/10], Loss: 0.0071
Epoch [6/10], Loss: 0.0069
Epoch [7/10], Loss: 0.0067
Epoch [8/10], Loss: 0.0066
Epoch [9/10], Loss: 0.0065
Epoch [10/10], Loss: 0.0064
```

```
In [ ]: plt.plot(range(1,len(training_loss_hid64)+1),training_loss_hid64,label="Hidden Size=64")
plt.plot(range(1,len(training_loss_hid128)+1),training_loss_hid128,label="Hidden Size=128")
plt.plot(range(1,len(training_loss_hid256)+1),training_loss_hid256,label="Hidden Size=256")
plt.legend()
plt.grid()
plt.title("Training Loss Vs Iterations")
plt.xlabel("Number of Iterations")
plt.ylabel("Training Loss")
plt.show()
```



```

In [ ]: model_hid64.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        outputs_hid64,_ = model_hid64(images.float())

model_hid128.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        outputs_hid128,_ = model_hid128(images.float())

model_hid256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        outputs_hid256,activations_hid256 = model_hid256(images.float())

plt.rcParams["figure.figsize"] = (15,6)
i=5
if i==5:
    fig, (ax1, ax2, ax3, ax4) = plt.subplots(1,4)

    ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_hid64[i].detach().numpy().reshape(28,28),cmap='gray')
    ax2.set_title('AE_hid64 Reconstructed Image')
    ax2.axis("off")


    ax3.imshow(outputs_hid128[i].detach().numpy().reshape(28,28),cmap='gray')
    ax3.set_title('AE_hid128 Reconstructed Image')
    ax3.axis("off")

    ax4.imshow(outputs_hid256[i].detach().numpy().reshape(28,28),cmap='gray')
    ax4.set_title('AE_hid256 Reconstructed Image')
    ax4.axis("off")

    print("Reconstruction Error in AE_hid64:",np.dot(((images[i].detach().numpy())/255
    print("Reconstruction Error in AE_hid128:",np.dot(((images[i].detach().numpy())/25
    print("Reconstruction Error in AE_hid256:",np.dot(((images[i].detach().numpy())/25

Reconstruction Error in AE_hid64: 22.084072923718914
Reconstruction Error in AE_hid128: 15.481209502901311
Reconstruction Error in AE_hid256: 12.704462646306368

```



```
In [ ]: #FASHION MNIST OUTPUT

# Load Fashion MNIST dataset
test_data_fashion = torchvision.datasets.FashionMNIST(root=".", train=False, transform=transforms.ToTensor())
fashion_image_sample = test_data_fashion.data[10]

# Set the models to evaluation mode and compute outputs
models = [model_hid64, model_hid128, model_hid256]
outputs = []
losses = []

# Criterion for calculating reconstruction loss
criterion = nn.MSELoss()

plt.rcParams["figure.figsize"] = (15, 6)

fig, (ax1, ax2, ax3, ax4) = plt.subplots(1, 4)

# Original Image
ax1.imshow(fashion_image_sample.detach().numpy().reshape(28, 28), cmap='gray')
ax1.set_title('Original Random Image')
ax1.axis("off")


for i, model in enumerate(models):
    model.eval()
    with torch.no_grad():
        images = fashion_image_sample.reshape(1, 28 * 28)
        output, _ = model(images.float())
        outputs.append(output)

        # Calculate reconstruction loss
        loss = criterion(output, images.float())
        losses.append(loss.item())

        # Plot reconstructed images
        ax = ax2 if i == 0 else (ax3 if i == 1 else ax4)
        ax.imshow(output.detach().numpy().reshape(28, 28), cmap='gray')
        ax.set_title(f'AE_hid{[64, 128, 256][i]} Reconstructed Image\nLoss: {losses[i]:.4f}')
        ax.axis("off")

plt.show()

# Print reconstruction losses
for i, loss in enumerate(losses):
    print(f'Reconstruction Loss for AE_hid{[64, 128, 256][i]}: {loss:.4f}')
```



Reconstruction Loss for AE_hid64: 6392.3477
Reconstruction Loss for AE_hid128: 5284.0664
Reconstruction Loss for AE_hid256: 4222.0620

```
In [ ]: # OUTPUTS FOR RANDOM IMAGE
torch.manual_seed(0)
random_image = torch.randint(low=0, high=255, size=(1, 28, 28))

# Initialize lists to store the reconstruction errors
reconstruction_errors = {}

# Evaluate model with hidden size 64
model_hid64.eval()
with torch.no_grad():
    images = random_image.reshape(1, 28 * 28)
    outputs_hid64, _ = model_hid64(images.float())
    reconstruction_errors['AE_hid64'] = np.dot(((images.numpy()/255) - (outputs_hid
                                                ((images.numpy()/255) - (outputs_hi

# Evaluate model with hidden size 128
model_hid128.eval()
with torch.no_grad():
    images = random_image.reshape(1, 28 * 28)
    outputs_hid128, _ = model_hid128(images.float())
    reconstruction_errors['AE_hid128'] = np.dot(((images.numpy()/255) - (outputs_hi
                                                ((images.numpy()/255) - (outputs_h

# Evaluate model with hidden size 256
model_hid256.eval()
with torch.no_grad():
    images = random_image.reshape(1, 28 * 28)
    outputs_hid256, _ = model_hid256(images.float())
    reconstruction_errors['AE_hid256'] = np.dot(((images.numpy()/255) - (outputs_hi
                                                ((images.numpy()/255) - (outputs_h

# Plotting
plt.rcParams["figure.figsize"] = (15, 6)
fig, (ax1, ax2, ax3, ax4) = plt.subplots(1, 4)


# Display the original random image
ax1.imshow(random_image.numpy().reshape(28, 28), cmap='gray')
ax1.set_title('Original Random Image')
ax1.axis("off")

# Display the reconstructed images and print reconstruction errors
ax2.imshow(outputs_hid64.detach().numpy().reshape(28, 28), cmap='gray')
ax2.set_title('AE_hid64 Reconstructed Image\nError: {:.4f}'.format(reconstruction_e
ax2.axis("off")

ax3.imshow(outputs_hid128.detach().numpy().reshape(28, 28), cmap='gray')
ax3.set_title('AE_hid128 Reconstructed Image\nError: {:.4f}'.format(reconstruction_
ax3.axis("off")

ax4.imshow(outputs_hid256.detach().numpy().reshape(28, 28), cmap='gray')
ax4.set_title('AE_hid256 Reconstructed Image\nError: {:.4f}'.format(reconstruction_
ax4.axis("off")

plt.show()
```

**Observations:**

- For an image sample from MNIST dataset, AE with hidden size 256 seems to reconstruct the image in a better manner. This was expected as more components would be able to reproduce the image in a better manner by capturing the minor details of the input image. This is also proved by the lesser reconstruction error by this model.
- For an image from Fashion MNIST dataset, the models try to reconstruct the image but the models with 64 and 128 hidden size could not reconstruct properly. On the other hand, 256 hidden size model does a better job than them but also seems to struggle in the reconstruction, though if the model was trained on fashion mnist dataset, it would've performed better.
- For a random image, the models perform as if they were trying to find out the digits, so the center pixels have some noisy output and the outer/edge pixels seem to be off (similar to the MNIST digit dataset, as our model has been trained on them.)

```
In [ ]: import torch
import matplotlib.pyplot as plt
import numpy as np

# Extract the weights from the first encoder layer
weights = model_hid64.encoder[0].weight.detach().numpy() # Adjust the layer index

# Normalize the weights to the range [0, 1] for visualization
weights_min = np.min(weights)
weights_max = np.max(weights)
weights_normalized = (weights - weights_min) / (weights_max - weights_min)

# Determine the number of filters (hidden nodes)
num_filters = weights_normalized.shape[0]
# Calculate the grid size for visualization
grid_size = int(np.ceil(np.sqrt(num_filters)))


# Set the figure size for the plots
plt.rcParams["figure.figsize"] = (12, 12)

# Create a grid to display the filters
fig, axes = plt.subplots(grid_size, grid_size)

for i in range(grid_size):
    for j in range(grid_size):
        index = i * grid_size + j
        if index < num_filters:
            ax = axes[i, j]
            ax.imshow(weights_normalized[index].reshape(28, 28), cmap='gray') # Re
            ax.axis('off')
        else:
            axes[i, j].axis('off') # Turn off unused axes

plt.suptitle('Learned Filters of Autoencoder (Hidden Size 64)', fontsize=16)
plt.show()
```


Learned Filters of Autoencoder (Hidden Size 64)



Edge Detection:

- Many filters highlight vertical and horizontal edges, which are essential for recognizing digit outlines. Digit Components:
- Some filters capture features specific to certain digits, such as the curvature of '0', '6', and '9'.

Shape and Structure:

- Filters often represent structural features, such as the loops in '8' or the straight lines in '1', aiding in digit classification.

Variability:

- Filters show variability across different digits, suggesting that the Autoencoder has learned to distinguish between shapes and styles of the numbers.

Hierarchical Learning:

- Lower layers may learn basic features (e.g., edges and corners), while higher layers capture more complex features (e.g., complete digits).

Noise Robustness:

- Some filters appear to generalize well to variations in handwriting styles, indicating the model's ability to handle noise.

Dimensionality Reduction:

- The learned filters can be seen as a compressed representation of the digits, facilitating efficient reconstruction and analysis of the digit images.

These observations suggest that the Autoencoder has effectively learned to represent the essential features of MNIST digits, highlighting its potential for applications in tasks like digit recognition and data compression.

Sparse Autoencoders

```
In [ ]: class AE3_SparseAutoencoder(nn.Module):
    def __init__(self):
        super(AE3_SparseAutoencoder, self).__init__()
        self.encoder = nn.Sequential(
            nn.Linear(784,1156),
            nn.ReLU())
        self.decoder = nn.Sequential(
            nn.Linear(1156,784),
            nn.ReLU())

    def forward(self,x):
        x=self.encoder(x)
        encoded_output=x
        l1_norm=torch.norm(x,p=1)
        x=self.decoder(x)
        return x,l1_norm,encoded_output
```

```
In [ ]: # Set lambda for L1 regularization
lambda_ = 1 * 1e-6
epochs= 10
# Initialize model, criterion, and optimizer
model_3_a = AE3_SparseAutoencoder()
criterion_3_a = nn.MSELoss()
optimizer_3_a = torch.optim.Adam(model_3_a.parameters(), lr=learning_rate)

# List to store training losses
training_loss_3_a = []

# Training loop
for epoch in range(epochs):
    epoch_loss = 0 # Variable to accumulate epoch loss
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1)

        # Forward pass
        outputs, l1_norm, _ = model_3_a(images)

        # Calculate loss with L1 regularization
        loss = criterion_3_a(outputs, images) + lambda_ * l1_norm

        # Backward pass and optimization
        optimizer_3_a.zero_grad()
        loss.backward()
        optimizer_3_a.step()

        # Accumulate loss for the epoch
        epoch_loss += loss.item()

    # Calculate average loss for the epoch
    avg_epoch_loss = epoch_loss / len(train_loader)
    training_loss_3_a.append(avg_epoch_loss)

    # Print epoch loss
    print(f"Epoch [{epoch + 1}/{epochs}] - Loss: {avg_epoch_loss:.4f} (including L1
```

```
Epoch [1/10] - Loss: 0.0295 (including L1 regularization)
Epoch [2/10] - Loss: 0.0150 (including L1 regularization)
Epoch [3/10] - Loss: 0.0128 (including L1 regularization)
Epoch [4/10] - Loss: 0.0116 (including L1 regularization)
Epoch [5/10] - Loss: 0.0108 (including L1 regularization)
Epoch [6/10] - Loss: 0.0102 (including L1 regularization)
Epoch [7/10] - Loss: 0.0097 (including L1 regularization)
Epoch [8/10] - Loss: 0.0093 (including L1 regularization)
Epoch [9/10] - Loss: 0.0089 (including L1 regularization)
Epoch [10/10] - Loss: 0.0086 (including L1 regularization)
```

```
In [ ]: # Set lambda for L1 regularization
lambda_ = 1 * 1e-5
epochs= 10
# Initialize model, criterion, and optimizer
model_3_b = AE3_SparseAutoencoder()
criterion_3_b = nn.MSELoss()
optimizer_3_b = torch.optim.Adam(model_3_b.parameters(), lr=learning_rate)

# List to store training losses
training_loss_3_b = []

# Training Loop
for epoch in range(epochs):
    epoch_loss = 0 # Variable to accumulate epoch loss
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1)

        # Forward pass
        outputs, l1_norm, _ = model_3_b(images)

        # Calculate loss with L1 regularization
        loss = criterion_3_b(outputs, images) + lambda_ * l1_norm

        # Backward pass and optimization
        optimizer_3_b.zero_grad()
        loss.backward()
        optimizer_3_b.step()

        # Accumulate loss for the epoch
        epoch_loss += loss.item()

    # Calculate average loss for the epoch
    avg_epoch_loss = epoch_loss / len(train_loader)
    training_loss_3_b.append(avg_epoch_loss)

    # Print epoch loss
    print(f"Epoch [{epoch + 1}/{epochs}] - Loss: {avg_epoch_loss:.4f} (including L1
```

```
Epoch [1/10] - Loss: 0.0969 (including L1 regularization)
Epoch [2/10] - Loss: 0.0812 (including L1 regularization)
Epoch [3/10] - Loss: 0.0781 (including L1 regularization)
Epoch [4/10] - Loss: 0.0775 (including L1 regularization)
Epoch [5/10] - Loss: 0.0775 (including L1 regularization)
Epoch [6/10] - Loss: 0.0775 (including L1 regularization)
Epoch [7/10] - Loss: 0.0775 (including L1 regularization)
Epoch [8/10] - Loss: 0.0775 (including L1 regularization)
Epoch [9/10] - Loss: 0.0775 (including L1 regularization)
Epoch [10/10] - Loss: 0.0775 (including L1 regularization)
```

```

In [ ]: # Set Lambda for L1 regularization
lambda_ = 1 * 1e-4
epochs= 10
# Initialize model, criterion, and optimizer
model_3_c = AE3_SparseAutoencoder()
criterion_3_c = nn.MSELoss()
optimizer_3_c = torch.optim.Adam(model_3_c.parameters(), lr=learning_rate)

# List to store training losses
training_loss_3_c = []

# Training loop
for epoch in range(epochs):
    epoch_loss = 0 # Variable to accumulate epoch loss
    for images, labels in train_loader:
        images = images.reshape(images.shape[0], -1)

        # Forward pass
        outputs, l1_norm, _ = model_3_c(images)

        # Calculate loss with L1 regularization
        loss = criterion_3_c(outputs, images) + lambda_ * l1_norm

        # Backward pass and optimization
        optimizer_3_c.zero_grad()
        loss.backward()
        optimizer_3_c.step()

        # Accumulate loss for the epoch
        epoch_loss += loss.item()

    # Calculate average loss for the epoch
    avg_epoch_loss = epoch_loss / len(train_loader)
    training_loss_3_c.append(avg_epoch_loss)

    # Print epoch loss
    print(f"Epoch [{epoch + 1}/{epochs}] - Loss: {avg_epoch_loss:.4f} (including L1

```

```


Epoch [1/10] - Loss: 0.1128 (including L1 regularization)
Epoch [2/10] - Loss: 0.0828 (including L1 regularization)
Epoch [3/10] - Loss: 0.0800 (including L1 regularization)
Epoch [4/10] - Loss: 0.0795 (including L1 regularization)
Epoch [5/10] - Loss: 0.0795 (including L1 regularization)
Epoch [6/10] - Loss: 0.0795 (including L1 regularization)
Epoch [7/10] - Loss: 0.0794 (including L1 regularization)
Epoch [8/10] - Loss: 0.0793 (including L1 regularization)
Epoch [9/10] - Loss: 0.0793 (including L1 regularization)
Epoch [10/10] - Loss: 0.0793 (including L1 regularization)

```

```

In [ ]: plt.plot(range(1,len(training_loss_3_a)+1),training_loss_3_a,label="λ=$1*10^{-6}$")
plt.plot(range(1,len(training_loss_3_a)+1),training_loss_3_b,label="λ=$1*10^{-5}$")
plt.plot(range(1,len(training_loss_3_a)+1),training_loss_3_c,label="λ=$1*10^{-4}$")
plt.legend()
plt.grid()
plt.title("Training Loss Vs Iterations")
plt.xlabel("Number of Iterations")
plt.ylabel("Training Loss")
plt.show()

```



```

In [ ]: model_3_a.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        outputs_3_a,_,activation_3_a = model_3_a(images.float())

model_3_b.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        outputs_3_b,_,activation_3_b = model_3_b(images.float())

model_3_c.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        outputs_3_c,_,activation_3_c = model_3_c(images.float())

plt.rcParams["figure.figsize"] = (15,6)

for i in range(10):
    fig, (ax1, ax2,ax3,ax4) = plt.subplots(1,4)

    ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_3_a[i].detach().numpy().reshape(28,28),cmap='gray')
    ax2.set_title('Sparse AE lambda=$1*10^{-6}$')
    ax2.axis("off")


    ax3.imshow(outputs_3_b[i].detach().numpy().reshape(28,28),cmap='gray')
    ax3.set_title('Sparse AE lambda=$1*10^{-5}$')
    ax3.axis("off")

    ax4.imshow(outputs_3_c[i].detach().numpy().reshape(28,28),cmap='gray')
    ax4.set_title('Sparse AE lambda=$1*10^{-4}$')
    ax4.axis("off")


    plt.show()

    print("Reconstruction Error in SparseAE  $\lambda = 1*10^{-6}$ :",np.dot(((images[i].detach().
    print("Reconstruction Error in SparseAE  $\lambda = 1*10^{-5}$ :",np.dot(((images[i].detach().
    print("Reconstruction Error in SparseAE  $\lambda = 1*10^{-4}$ :",np.dot(((images[i].detach().
    print("


```




Reconstruction Error in SparseAE $\lambda = 1*1e-6$: 3.562504673493041
 Reconstruction Error in SparseAE $\lambda = 1*1e-5$: 47.76665086547913
 Reconstruction Error in SparseAE $\lambda = 1*1e-4$: 47.76881409034138




Reconstruction Error in SparseAE $\lambda = 1*1e-6$: 13.059560942836569
 Reconstruction Error in SparseAE $\lambda = 1*1e-5$: 114.2811292577155
 Reconstruction Error in SparseAE $\lambda = 1*1e-4$: 114.30503260973435




Reconstruction Error in SparseAE $\lambda = 1*1e-6$: 11.979831259337601
 Reconstruction Error in SparseAE $\lambda = 1*1e-5$: 119.60605315038694
 Reconstruction Error in SparseAE $\lambda = 1*1e-4$: 119.609050199548




Reconstruction Error in SparseAE $\lambda = 1*1e-6$: 6.84872849367996
 Reconstruction Error in SparseAE $\lambda = 1*1e-5$: 79.08692220978871
 Reconstruction Error in SparseAE $\lambda = 1*1e-4$: 79.10138781248455




Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-6}$: 6.414019453175997
 Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-5}$: 55.07611049870571
 Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-4}$: 55.08305633264402




Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-6}$: 13.765464797562657
 Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-5}$: 131.18994794171863
 Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-4}$: 131.21599974547468



Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-6}$: 7.269245388821619
 Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-5}$: 73.52382924375961
 Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-4}$: 73.54789156876805




Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-6}$: 12.288598237094437
 Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-5}$: 102.02350773752107
 Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-4}$: 102.02646721570278



Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-6}$: 9.012632636878312

Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-5}$: 70.64328764713406

Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-4}$: 70.66243320756527



Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-6}$: 12.760878562161171

Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-5}$: 115.58458519082258

Reconstruction Error in SparseAE $\lambda = 1 \cdot 10^{-4}$: 115.59485376822968

```
In [ ]: ##VISUALISING ACTIVATIONS

plt.rcParams["figure.figsize"] = (15,6)
for i in range(10):
    fig, (ax1, ax2, ax3, ax4, ax5) = plt.subplots(1,5)
    ax1.imshow(images[i].detach().numpy().reshape(28,28), cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")


    ax2.imshow(np.array(activations_hid256.detach().numpy())[i].reshape(int(np.sqrt(256))))
    ax2.set_title('StandardAE')
    ax2.axis("off")


    ax3.imshow(np.array(activation_3_a.detach().numpy())[i].reshape(int(np.sqrt(1156))))
    ax3.set_title('SparseAE  $\lambda=1 \cdot 10^{-6}$ ')
    ax3.axis("off")


    ax4.imshow(np.array(activation_3_b.detach().numpy())[i].reshape(int(np.sqrt(1156))))
    ax4.set_title('SparseAE  $\lambda=1 \cdot 10^{-5}$ ')
    ax4.axis("off")

    ax5.imshow(np.array(activation_3_c.detach().numpy())[i].reshape(int(np.sqrt(1156))))
    ax5.set_title('SparseAE  $\lambda=1 \cdot 10^{-4}$ ')
    ax5.axis("off")

plt.show()
```







```
In [ ]: ##VISUALISING ENCODER WEIGHTS

plt.rcParams["figure.figsize"] = (15,6)


fig, (ax1, ax2,ax3,ax4) = plt.subplots(1,4)
ax1.imshow(model_hid256.encoder[0].weight.detach().numpy()[0].reshape(28,28), cmap='grayscale')
ax1.set_title('StandradAE')
ax1.axis("off")

ax2.imshow(model_3_a.encoder[0].weight.detach().numpy()[0].reshape(28,28), cmap='grayscale')
ax2.set_title('SparseAE λ=$1*10^{-6}$')
ax2.axis("off")

ax3.imshow(model_3_b.encoder[0].weight.detach().numpy()[0].reshape(28,28), cmap='grayscale')
ax3.set_title('SparseAE λ=$1*10^{-5}$')
ax3.axis("off")

ax4.imshow(model_3_c.encoder[0].weight.detach().numpy()[0].reshape(28,28), cmap='grayscale')
ax4.set_title('SparseAE λ=$1*10^{-4}$')
ax4.axis("off")

plt.show()
```



```
In [ ]: ##VISUALISING ENCODER WEIGHTS

plt.rcParams["figure.figsize"] = (15,6)


fig, (ax1, ax2, ax3, ax4) = plt.subplots(1,4)
ax1.imshow(model_hid256.encoder[0].weight.detach().numpy()[6].reshape(28,28), cmap='grayscale')
ax1.set_title('StandradAE')
ax1.axis("off")

ax2.imshow(model_3_a.encoder[0].weight.detach().numpy()[6].reshape(28,28), cmap='grayscale')
ax2.set_title('SparseAE  $\lambda=1*10^{-6}$ ')
ax2.axis("off")

ax3.imshow(model_3_b.encoder[0].weight.detach().numpy()[6].reshape(28,28), cmap='grayscale')
ax3.set_title('SparseAE  $\lambda=1*10^{-5}$ ')
ax3.axis("off")

ax4.imshow(model_3_c.encoder[0].weight.detach().numpy()[6].reshape(28,28), cmap='grayscale')
ax4.set_title('SparseAE  $\lambda=1*10^{-4}$ ')
ax4.axis("off")

plt.show()
```



Observations:

- As lambda(sparsity factor) is decreased, reconstruction error goes down.
- If the lambda is continually decreased beyond a point then the model starts to overfit data. since the model is overcomplete, we need to iterate over the optimal lambda parameter to get desired output.
- as we decrease lambda(sparsity parameter) more neurons get activated. As can be observed from activation visualisations. However, in vanilla AE all the neurons are activated.
- Visualisation of encoder weights is difficult as nothing could be inferred from the images.

Denoising Autoencoders

```
In [ ]: class AE4_DenoisingAutoencoder(nn.Module):
        def __init__(self):
            super(AE4_DenoisingAutoencoder, self).__init__()
            self.encoder = nn.Sequential(
                nn.Linear(784,256),
                nn.ReLU())
            self.decoder =nn.Sequential(
                nn.Linear(256,784),
                nn.ReLU())

        def forward(self,x):
            x=self.encoder(x)
            x=self.decoder(x)
            return x
```

```
In [ ]: def add_noise(img, noise_val):
        noise = torch.randn(img.size())*noise_val
        noisy_img = img + noise
        return noisy_img
```

```
In [ ]: #HYPERPARAMETERS
        learning_rate=0.0008
        epochs=10
        batch_size= 256
```

```
In [ ]: model_4_a=AE4_DenoisingAutoencoder()
        criterion_4_a=nn.MSELoss()
        optimizer_4_a = torch.optim.Adam(model_4_a.parameters(),lr=learning_rate)

        training_loss_4_a=[]

        for epoch in range(epochs):
            for images,labels in train_loader:
                images=images.reshape(images.shape[0],-1)
                noisy_images=add_noise(images,0.3)
                outputs=model_4_a(noisy_images)
                loss=criterion_4_a(outputs,images)
                training_loss_4_a.append(loss.item())

                optimizer_4_a.zero_grad()
                loss.backward()
                optimizer_4_a.step()

            print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

```
In [ ]: model_4_b=AE4_DenoisingAutoencoder()
criterion_4_b=nn.MSELoss()
optimizer_4_b = torch.optim.Adam(model_4_b.parameters(),lr=learning_rate)

training_loss_4_b=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        images=images.reshape(images.shape[0],-1)
        noisy_images=add_noise(images,0.5)
        outputs=model_4_b(noisy_images)
        loss=criterion_4_b(outputs,images)
        training_loss_4_b.append(loss.item())

        optimizer_4_b.zero_grad()
        loss.backward()
        optimizer_4_b.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

```
In [ ]: model_4_c=AE4_DenoisingAutoencoder()
criterion_4_c=nn.MSELoss()
optimizer_4_c = torch.optim.Adam(model_4_c.parameters(),lr=learning_rate)

training_loss_4_c=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        images=images.reshape(images.shape[0],-1)
        noisy_images=add_noise(images,0.8)
        outputs=model_4_c(noisy_images)
        loss=criterion_4_c(outputs,images)
        training_loss_4_c.append(loss.item())

        optimizer_4_c.zero_grad()
        loss.backward()
        optimizer_4_c.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

```
In [ ]: model_4_d=AE4_DenoisingAutoencoder()
        criterion_4_d=nn.MSELoss()
        optimizer_4_d = torch.optim.Adam(model_4_d.parameters(),lr=learning_rate)

        training_loss_4_d=[]

        for epoch in range(epochs):
            for images,labels in train_loader:
                images=images.reshape(images.shape[0],-1)
                noisy_images=add_noise(images,0.9)
                outputs=model_4_d(noisy_images)
                loss=criterion_4_d(outputs,images)
                training_loss_4_d.append(loss.item())

                optimizer_4_d.zero_grad()
                loss.backward()
                optimizer_4_d.step()

            print("Epoch", "[",epoch+1,"/",epochs,"]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

```
In [ ]: plt.rcParams["figure.figsize"] = (15,6)

        plt.plot(range(1,len(training_loss_4_a)+1),training_loss_4_a,label="noise factor =0")
        plt.plot(range(1,len(training_loss_4_a)+1),training_loss_4_b,label="noise factor =0")
        plt.plot(range(1,len(training_loss_4_a)+1),training_loss_4_c,label="noise factor =0")
        plt.plot(range(1,len(training_loss_4_a)+1),training_loss_4_d,label="noise factor =0")
        plt.legend()
        plt.grid()
        plt.title("Training Loss Vs Iterations")
        plt.xlabel("Number of Iterations")
        plt.ylabel("Training Loss")
        plt.show()
```




```

In [ ]: #INPUT WITH NOISE FACTOR=0.3,0.5,0.8,0.9 GIVEN TO STANDARD AE HIDLAYER=256
model_hid256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        noisy_images=add_noise(images,0.3)
        outputs_hid256_03,activations_hid256 = model_hid256(noisy_images.float())

model_hid256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images=add_noise(images,0.5)
        outputs_hid256_05,activations_hid256 = model_hid256(noisy_images.float())

model_hid256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images=add_noise(images,0.8)
        outputs_hid256_08,activations_hid256 = model_hid256(noisy_images.float())

model_hid256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images=add_noise(images,0.9)
        outputs_hid256_09,activations_hid256 = model_hid256(noisy_images.float())

plt.rcParams["figure.figsize"] = (15,6)
i=5
fig, (ax1, ax2,ax3,ax4,ax5) = plt.subplots(1,5)

ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
ax1.set_title('Original Image')
ax1.axis("off")

ax2.imshow(outputs_hid256_03[i].detach().numpy().reshape(28,28),cmap='gray')
ax2.set_title('AE_Reconst.Image_nf=0.3')
ax2.axis("off")

ax3.imshow(outputs_hid256_05[i].detach().numpy().reshape(28,28),cmap='gray')
ax3.set_title('AE_Reconst.Image_nf=0.5')
ax3.axis("off")


ax4.imshow(outputs_hid256_08[i].detach().numpy().reshape(28,28),cmap='gray')
ax4.set_title('AE_Reconst.Image_nf=0.8')
ax4.axis("off")

ax5.imshow(outputs_hid256_09[i].detach().numpy().reshape(28,28),cmap='gray')
ax5.set_title('AE_Reconst.Image_nf=0.9')
ax5.axis("off")

print("Reconstruction Error in StandardAE with noise factor = 0.3 :",np.dot(((image
print("Reconstruction Error in StandardAE with noise factor = 0.5 :",np.dot(((image
print("Reconstruction Error in StandardAE with noise factor = 0.8 :",np.dot(((image
print("Reconstruction Error in StandardAE with noise factor = 0.9 :",np.dot(((image

```

Reconstruction Error in StandardAE with noise factor = 0.3 : 12.708732499861467
Reconstruction Error in StandardAE with noise factor = 0.5 : 12.689272416002737
Reconstruction Error in StandardAE with noise factor = 0.8 : 12.704175335761287
Reconstruction Error in StandardAE with noise factor = 0.9 : 12.711788004726792



```

In [ ]: model_4_a.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images,0.3)
        outputs_4_a = model_4_a(noisy_images.float())

model_4_b.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images,0.5)
        outputs_4_b = model_4_b(noisy_images.float())

model_4_c.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images,0.8)
        outputs_4_c = model_4_c(noisy_images.float())

model_4_d.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images,0.9)
        outputs_4_d = model_4_d(noisy_images.float())

plt.rcParams["figure.figsize"] = (15,6)
i=5
if i==5:
    fig, (ax1, ax2,ax3,ax4,ax5) = plt.subplots(1,5)

    ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_4_a[i].detach().numpy().reshape(28,28),cmap='gray')
    ax2.set_title('Reconstruction: nf=0.3')
    ax2.axis("off")

    ax3.imshow(outputs_4_b[i].detach().numpy().reshape(28,28),cmap='gray')
    ax3.set_title('Reconstruction: nf=0.5')
    ax3.axis("off")

    ax4.imshow(outputs_4_c[i].detach().numpy().reshape(28,28),cmap='gray')
    ax4.set_title('Reconstruction: nf=0.8')
    ax4.axis("off")


    ax5.imshow(outputs_4_d[i].detach().numpy().reshape(28,28),cmap='gray')
    ax5.set_title('Reconstruction: nf=0.9')
    ax5.axis("off")

plt.show()

print("Reconstruction Error in DenoisingAE with noise factor = 0.3 :",np.dot(((im
print("Reconstruction Error in DenoisingAE with noise factor = 0.5 :",np.dot(((im
print("Reconstruction Error in DenoisingAE with noise factor = 0.8 :",np.dot(((im
print("Reconstruction Error in DenoisingAE with noise factor = 0.9 :",np.dot(((im

```

```
print("Reconstruction Error in DenoisingAE with noise factor = 0.9 : ", np.dot((im  
print("_____
```



```
Reconstruction Error in DenoisingAE with noise factor = 0.3 : 11.899540434699826  
Reconstruction Error in DenoisingAE with noise factor = 0.5 : 11.693956936036919  
Reconstruction Error in DenoisingAE with noise factor = 0.8 : 20.232822123745233  
Reconstruction Error in DenoisingAE with noise factor = 0.9 : 21.920191710208133
```

```

In [ ]: #DENOISING ENCODER TRAINED ON NOISE LEVEL=0.3 AND TESTED RECONSTRUCTION FOR NOISE L
model_4_a.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images,0.3)
        outputs_4_a = model_4_a(noisy_images.float())
        noisy_images = add_noise(images,0.5)
        outputs_4_b = model_4_a(noisy_images.float())
        noisy_images = add_noise(images,0.8)
        outputs_4_c = model_4_a(noisy_images.float())
        noisy_images = add_noise(images,0.9)
        outputs_4_d = model_4_a(noisy_images.float())
i=5
if i==5:
    fig, (ax1, ax2,ax3,ax4,ax5) = plt.subplots(1,5)

    ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_4_a[i].detach().numpy().reshape(28,28),cmap='gray')
    ax2.set_title('Reconstruction: nf=0.3')
    ax2.axis("off")

    ax3.imshow(outputs_4_b[i].detach().numpy().reshape(28,28),cmap='gray')
    ax3.set_title('Reconstruction: nf=0.5')
    ax3.axis("off")


    ax4.imshow(outputs_4_c[i].detach().numpy().reshape(28,28),cmap='gray')
    ax4.set_title('Reconstruction: nf=0.8')
    ax4.axis("off")

    ax5.imshow(outputs_4_d[i].detach().numpy().reshape(28,28),cmap='gray')
    ax5.set_title('Reconstruction: nf=0.9')
    ax5.axis("off")

    plt.show()

    print("Reconstruction Error in DenoisingAE with noise factor = 0.3 :",np.dot(((im
    print("Reconstruction Error in DenoisingAE with noise factor = 0.5 :",np.dot(((im
    print("Reconstruction Error in DenoisingAE with noise factor = 0.8 :",np.dot(((im
    print("Reconstruction Error in DenoisingAE with noise factor = 0.9 :",np.dot(((im
    print("

```



```

Reconstruction Error in DenoisingAE with noise factor = 0.3 : 11.894745455874139
Reconstruction Error in DenoisingAE with noise factor = 0.5 : 11.898799213156467
Reconstruction Error in DenoisingAE with noise factor = 0.8 : 11.89442192037134
Reconstruction Error in DenoisingAE with noise factor = 0.9 : 11.895282881344349

```

```
In [ ]: ##VISUALISING ENCODER WEIGHTS

plt.rcParams["figure.figsize"] = (15,6)

fig, (ax1, ax2, ax3, ax4, ax5) = plt.subplots(1, 5)
ax1.imshow(model_hid256.encoder[0].weight.detach().numpy()[0].reshape(28,28), cmap='grayscale')
ax1.set_title('StandardAE')
ax1.axis("off")


ax2.imshow(model_4_a.encoder[0].weight.detach().numpy()[0].reshape(28,28), cmap='grayscale')
ax2.set_title('DenoisingAE nf=0.3')
ax2.axis("off")

ax3.imshow(model_4_b.encoder[0].weight.detach().numpy()[0].reshape(28,28), cmap='grayscale')
ax3.set_title('DenoisingAE nf=0.5')
ax3.axis("off")

ax4.imshow(model_4_c.encoder[0].weight.detach().numpy()[0].reshape(28,28), cmap='grayscale')
ax4.set_title('DenoisingAE nf=0.8')
ax4.axis("off")

ax5.imshow(model_4_d.encoder[0].weight.detach().numpy()[0].reshape(28,28), cmap='grayscale')
ax5.set_title('DenoisingAE nf=0.9')
ax5.axis("off")

plt.show()
```



Observations:

- Standard autoencoder when given corrupted input has reconstruction error more than Denoising Autoencoder trained on noise level 0.3.
- Though not very clearly visible but encoder weights of denoising AE have some penstroke detector type visualisation.

Manifold Learning

Impact of adding random noise to input images

```
In [8]: # Function to add random noise to the input images
def add_noise_to_input(images, noise_factor=0.5):
    noisy_images = images + noise_factor * torch.randn_like(images)
    noisy_images = torch.clip(noisy_images, 0., 1.) # Ensure pixel values remain i
    return noisy_images

# Get sample images from MNIST test set
sample_images, _ = next(iter(test_loader))
sample_images = sample_images[:10] # Take the first 10 images
sample_images_flat = sample_images.view(sample_images.size(0), -1) # Flatten the i


# Add random noise to input images
noisy_images = add_noise_to_input(sample_images_flat)

# Visualize the original and noisy images
plt.figure(figsize=(20, 5))

for i in range(10): # Loop through the 10 images
    # Plot original image
    plt.subplot(2, 10, i + 1)
    plt.imshow(sample_images[i].view(28, 28).cpu().numpy(), cmap='gray')
    plt.title(f'Original')
    plt.axis('off')

    # Plot noisy image
    plt.subplot(2, 10, i + 11) # Shift to the next row
    plt.imshow(noisy_images[i].view(28, 28).cpu().numpy(), cmap='gray')
    plt.title(f'Noisy')
    plt.axis('off')

plt.show()
```



Observation:

- Random noise disrupts the structured patterns that represent digits.
- The high-dimensional input space for MNIST (784 dimensions) requires specific arrangements of pixels to represent valid digits, and random changes break this structure, making it unlikely that the resulting noisy image still represents a valid digit.
- The noise introduces high-contrast pixels, increasing visual complexity and blurring the boundaries between the digits and the background, which challenges the isolation of original digit features.


```
In [9]: class Autoencoder(nn.Module):
        def __init__(self):
            super(Autoencoder, self).__init__()
            self.encoder = nn.Sequential(
                nn.Linear(784, 64),
                nn.ReLU(),
                nn.Linear(64, 8), # Bottleneck
                nn.ReLU()
            )
            self.decoder = nn.Sequential(
                nn.Linear(8, 64),
                nn.ReLU(),
                nn.Linear(64, 784),
                nn.Sigmoid() # Using Sigmoid to match output range [0, 1]
            )

        def forward(self, x):
            x = self.encoder(x)
            x = self.decoder(x)
            return x
```

```
In [10]: # Initialize the model, Loss function, and optimizer
model = Autoencoder()
criterion = nn.MSELoss()
criterion = nn.MSELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

# Training Loop
epochs = 10
for epoch in range(epochs):
    for images, _ in train_loader:
        images = images.view(images.size(0), -1) # Flatten the images
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, images)
        loss.backward()
        optimizer.step()

    print(f'Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}')
```

```
Epoch [1/10], Loss: 0.0559
Epoch [2/10], Loss: 0.0509
Epoch [3/10], Loss: 0.0491
Epoch [4/10], Loss: 0.0410
Epoch [5/10], Loss: 0.0376
Epoch [6/10], Loss: 0.0386
Epoch [7/10], Loss: 0.0349
Epoch [8/10], Loss: 0.0333
Epoch [9/10], Loss: 0.0351
Epoch [10/10], Loss: 0.0344
```

```
In [11]: # Function to add noise to the representation
def add_noise(representation, noise_factor=0.5):
    noise = torch.randn_like(representation) * noise_factor
    return representation + noise

# Get two sample images from the test set after the model has converged
sample_images, _ = next(iter(test_loader))

# Ensure we only use 2 images from the batch
sample_images = sample_images[:2]
sample_images_flat = sample_images.view(2, -1) # Flatten to shape (2, 784) for MNI

# Pass the sample images through the trained autoencoder
with torch.no_grad():
    original_reconstructions = model(sample_images_flat) # Forward pass through au

# Encode the sample images (latent representation)
with torch.no_grad():
    encoded_representations = model.encoder(sample_images_flat)

# Add noise to the latent representation
noisy_representations = add_noise(encoded_representations)

# Decode the noisy representations (reconstruct from noisy latent space)
with torch.no_grad():
    noisy_reconstructions = model.decoder(noisy_representations)

# Visualize the original images, reconstructed images, and noisy reconstructed images
plt.figure(figsize=(15, 9))


for i in range(2): # Loop through the 2 images
    # Original Image
    plt.subplot(2, 3, i*3 + 1)
    plt.imshow(sample_images[i].squeeze(), cmap='gray')
    plt.title(f'Original Image {i+1}')
    plt.axis('off')

    # Reconstructed Image (without noise)
    plt.subplot(2, 3, i*3 + 2)
    plt.imshow(original_reconstructions[i].view(28, 28).squeeze(), cmap='gray')
    plt.title(f'Reconstructed Image {i+1}')
    plt.axis('off')


    # Noisy Reconstructed Image
    plt.subplot(2, 3, i*3 + 3)
    plt.imshow(noisy_reconstructions[i].view(28, 28).squeeze(), cmap='gray')
    plt.title(f'Noisy Reconstruction {i+1}')
    plt.axis('off')

plt.show()
```


Original Image 1




Reconstructed Image 1




Noisy Reconstruction 1




Original Image 2



Reconstructed Image 2



Noisy Reconstruction 2



```
In [12]: # Function to add noise to the representation
def add_noise(representation, noise_factor=1.0):
    noise = torch.randn_like(representation) * noise_factor
    return representation + noise

# Get two sample images from the test set after the model has converged
sample_images, _ = next(iter(test_loader))

# Ensure we only use 2 images from the batch
sample_images = sample_images[:2]
sample_images_flat = sample_images.view(2, -1) # Flatten to shape (2, 784) for MNI

# Pass the sample images through the trained autoencoder
with torch.no_grad():
    original_reconstructions = model(sample_images_flat) # Forward pass through au

# Encode the sample images (latent representation)
with torch.no_grad():
    encoded_representations = model.encoder(sample_images_flat)

# Add noise to the latent representation
noisy_representations = add_noise(encoded_representations)

# Decode the noisy representations (reconstruct from noisy latent space)
with torch.no_grad():
    noisy_reconstructions = model.decoder(noisy_representations)

# Visualize the original images, reconstructed images, and noisy reconstructed images
plt.figure(figsize=(15, 9))


for i in range(2): # Loop through the 2 images
    # Original Image
    plt.subplot(2, 3, i*3 + 1)
    plt.imshow(sample_images[i].squeeze(), cmap='gray')
    plt.title(f'Original Image {i+1}')
    plt.axis('off')

    # Reconstructed Image (without noise)
    plt.subplot(2, 3, i*3 + 2)
    plt.imshow(original_reconstructions[i].view(28, 28).squeeze(), cmap='gray')
    plt.title(f'Reconstructed Image {i+1}')
    plt.axis('off')


    # Noisy Reconstructed Image
    plt.subplot(2, 3, i*3 + 3)
    plt.imshow(noisy_reconstructions[i].view(28, 28).squeeze(), cmap='gray')
    plt.title(f'Noisy Reconstruction {i+1}')
    plt.axis('off')

plt.show()
```


Original Image 1




Reconstructed Image 1




Noisy Reconstruction 1




Original Image 2



Reconstructed Image 2



Noisy Reconstruction 2



```
In [13]: # Function to add noise to the representation
def add_noise(representation, noise_factor=3.0):
    noise = torch.randn_like(representation) * noise_factor
    return representation + noise

# Get two sample images from the test set after the model has converged
sample_images, _ = next(iter(test_loader))

# Ensure we only use 2 images from the batch
sample_images = sample_images[:2]
sample_images_flat = sample_images.view(2, -1) # Flatten to shape (2, 784) for MNI

# Pass the sample images through the trained autoencoder
with torch.no_grad():
    original_reconstructions = model(sample_images_flat) # Forward pass through au

# Encode the sample images (latent representation)
with torch.no_grad():
    encoded_representations = model.encoder(sample_images_flat)

# Add noise to the latent representation
noisy_representations = add_noise(encoded_representations)

# Decode the noisy representations (reconstruct from noisy latent space)
with torch.no_grad():
    noisy_reconstructions = model.decoder(noisy_representations)

# Visualize the original images, reconstructed images, and noisy reconstructed images
plt.figure(figsize=(15, 9))


for i in range(2): # Loop through the 2 images
    # Original Image
    plt.subplot(2, 3, i*3 + 1)
    plt.imshow(sample_images[i].squeeze(), cmap='gray')
    plt.title(f'Original Image {i+1}')
    plt.axis('off')

    # Reconstructed Image (without noise)
    plt.subplot(2, 3, i*3 + 2)
    plt.imshow(original_reconstructions[i].view(28, 28).squeeze(), cmap='gray')
    plt.title(f'Reconstructed Image {i+1}')
    plt.axis('off')


    # Noisy Reconstructed Image
    plt.subplot(2, 3, i*3 + 3)
    plt.imshow(noisy_reconstructions[i].view(28, 28).squeeze(), cmap='gray')
    plt.title(f'Noisy Reconstruction {i+1}')
    plt.axis('off')

plt.show()
```


Original Image 1



Reconstructed Image 1




Noisy Reconstruction 1




Original Image 2



Reconstructed Image 2



Noisy Reconstruction 2



```
In [14]: # Function to add noise to the representation
def add_noise(representation, noise_factor=5.0):
    noise = torch.randn_like(representation) * noise_factor
    return representation + noise

# Get two sample images from the test set after the model has converged
sample_images, _ = next(iter(test_loader))

# Ensure we only use 2 images from the batch
sample_images = sample_images[:2]
sample_images_flat = sample_images.view(2, -1) # Flatten to shape (2, 784) for MNI

# Pass the sample images through the trained autoencoder
with torch.no_grad():
    original_reconstructions = model(sample_images_flat) # Forward pass through au

# Encode the sample images (latent representation)
with torch.no_grad():
    encoded_representations = model.encoder(sample_images_flat)

# Add noise to the latent representation
noisy_representations = add_noise(encoded_representations)

# Decode the noisy representations (reconstruct from noisy latent space)
with torch.no_grad():
    noisy_reconstructions = model.decoder(noisy_representations)

# Visualize the original images, reconstructed images, and noisy reconstructed images
plt.figure(figsize=(15, 9))


for i in range(2): # Loop through the 2 images
    # Original Image
    plt.subplot(2, 3, i*3 + 1)
    plt.imshow(sample_images[i].squeeze(), cmap='gray')
    plt.title(f'Original Image {i+1}')
    plt.axis('off')

    # Reconstructed Image (without noise)
    plt.subplot(2, 3, i*3 + 2)
    plt.imshow(original_reconstructions[i].view(28, 28).squeeze(), cmap='gray')
    plt.title(f'Reconstructed Image {i+1}')
    plt.axis('off')


    # Noisy Reconstructed Image
    plt.subplot(2, 3, i*3 + 3)
    plt.imshow(noisy_reconstructions[i].view(28, 28).squeeze(), cmap='gray')
    plt.title(f'Noisy Reconstruction {i+1}')
    plt.axis('off')

plt.show()
```



Original Image 1




Reconstructed Image 1




Noisy Reconstruction 1




Original Image 2



Reconstructed Image 2



Noisy Reconstruction 2



```
In [18]: # Function to add noise to the representation
def add_noise(representation, noise_factor=7.0):
    noise = torch.randn_like(representation) * noise_factor
    return representation + noise

# Get two sample images from the test set after the model has converged
sample_images, _ = next(iter(test_loader))

# Ensure we only use 2 images from the batch
sample_images = sample_images[:2]
sample_images_flat = sample_images.view(2, -1) # Flatten to shape (2, 784) for MNI

# Pass the sample images through the trained autoencoder
with torch.no_grad():
    original_reconstructions = model(sample_images_flat) # Forward pass through au

# Encode the sample images (latent representation)
with torch.no_grad():
    encoded_representations = model.encoder(sample_images_flat)

# Add noise to the latent representation
noisy_representations = add_noise(encoded_representations)

# Decode the noisy representations (reconstruct from noisy latent space)
with torch.no_grad():
    noisy_reconstructions = model.decoder(noisy_representations)

# Visualize the original images, reconstructed images, and noisy reconstructed images
plt.figure(figsize=(15, 9))


for i in range(2): # Loop through the 2 images
    # Original Image
    plt.subplot(2, 3, i*3 + 1)
    plt.imshow(sample_images[i].squeeze(), cmap='gray')
    plt.title(f'Original Image {i+1}')
    plt.axis('off')

    # Reconstructed Image (without noise)
    plt.subplot(2, 3, i*3 + 2)
    plt.imshow(original_reconstructions[i].view(28, 28).squeeze(), cmap='gray')
    plt.title(f'Reconstructed Image {i+1}')
    plt.axis('off')


    # Noisy Reconstructed Image
    plt.subplot(2, 3, i*3 + 3)
    plt.imshow(noisy_reconstructions[i].view(28, 28).squeeze(), cmap='gray')
    plt.title(f'Noisy Reconstruction {i+1}')
    plt.axis('off')

plt.show()
```


Original Image 1



Reconstructed Image 1




Noisy Reconstruction 1




Original Image 2



Reconstructed Image 2



Noisy Reconstruction 2



```
In [24]: # Function to add noise to the representation
def add_noise(representation, noise_factor=10.0):
    noise = torch.randn_like(representation) * noise_factor
    return representation + noise

# Get two sample images from the test set after the model has converged
sample_images, _ = next(iter(test_loader))

# Ensure we only use 2 images from the batch
sample_images = sample_images[:2]
sample_images_flat = sample_images.view(2, -1) # Flatten to shape (2, 784) for MNI

# Pass the sample images through the trained autoencoder
with torch.no_grad():
    original_reconstructions = model(sample_images_flat) # Forward pass through au

# Encode the sample images (latent representation)
with torch.no_grad():
    encoded_representations = model.encoder(sample_images_flat)

# Add noise to the latent representation
noisy_representations = add_noise(encoded_representations)

# Decode the noisy representations (reconstruct from noisy latent space)
with torch.no_grad():
    noisy_reconstructions = model.decoder(noisy_representations)


# Visualize the original images, reconstructed images, and noisy reconstructed images
plt.figure(figsize=(15, 9))

for i in range(2): # Loop through the 2 images
    # Original Image
    plt.subplot(2, 3, i*3 + 1)
    plt.imshow(sample_images[i].squeeze(), cmap='gray')
    plt.title(f'Original Image {i+1}')
    plt.axis('off')

    # Reconstructed Image (without noise)
    plt.subplot(2, 3, i*3 + 2)
    plt.imshow(original_reconstructions[i].view(28, 28).squeeze(), cmap='gray')
    plt.title(f'Reconstructed Image {i+1}')
    plt.axis('off')

    # Noisy Reconstructed Image
    plt.subplot(2, 3, i*3 + 3)
    plt.imshow(noisy_reconstructions[i].view(28, 28).squeeze(), cmap='gray')
    plt.title(f'Noisy Reconstruction {i+1}')
    plt.axis('off')

plt.show()
```



Observations:

- As the noise level increases, the quality of the reconstructed images deteriorates significantly. At lower noise levels (0.5-2.0), the autoencoder manages to reconstruct recognizable digits. However, beyond a noise level of 5.0, the reconstructions become increasingly distorted and less identifiable.
- Higher noise levels result in a more pronounced loss of detail in the reconstructed images. For instance, at a noise level of 5.0, some features, such as edges and curves of digits, become blurred, while at noise levels above 7.0, the digits appear almost as random patterns, making it difficult to discern the original digit.
- With higher noise levels (around 8.0-10.0), the outputs exhibit significant variability despite being generated from similar latent representations. This suggests that the autoencoder struggles to map the noisy latent space back to a coherent output, leading to a broader range of reconstruction outcomes that do not correspond to the original inputs.

Convolutional Autoencoders

```

In [ ]: class AE5_ConvAE_with_unpooling(nn.Module): #define unpooling outside the decoder a

def __init__(self): #class constructor
    super(AE5_ConvAE_with_unpooling,self).__init__() #calls the parent construc

    #initializing the encoder module
    self.encoder_conv1 = nn.Sequential(
        nn.Conv2d(1,8, kernel_size = 3, stride = 1,padding= 1),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
    ) # 28x28x1 to 14x14x8
    self.encoder_conv2 = nn.Sequential(
        nn.Conv2d(8,16, kernel_size = 3, stride = 1,padding= 1),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
    ) #14x14x8 to 7x7x16
    self.encoder_conv3 = nn.Sequential(
        nn.Conv2d(16,16, kernel_size = 3, stride = 1,padding= 1),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
    ) #7x7x16 to 3x3x16

    #initializing the decoder module
    self.decoder_conv1 = nn.Sequential(nn.Identity()) #7x7x16 to 7x7x16
    self.decoder_conv2 = nn.Sequential(
        nn.Conv2d(16,8, kernel_size = 3, stride = 1,padding= 1),
        nn.ReLU()
    ) #14x14x16 to 14x14x8
    self.decoder_conv3 = nn.Sequential(
        nn.Conv2d(8,1, kernel_size = 3, stride = 1,padding= 1),
        nn.ReLU()
    ) #28x28x8 to 28x28x1

    #defining the unpooling operation
    self.unpool = nn.MaxUnpool2d(kernel_size = (2,2))

def forward(self,x): #defines the forward pass and also the structure of the ne

    encoded_input,indices1 = self.encoder_conv1(x.float()) # 28x28x1 to 14x14x8
    encoded_input,indices2 = self.encoder_conv2(encoded_input) #14x14x8 to 7x7x16
    encoded_input,indices3 = self.encoder_conv3(encoded_input) #7x7x16 to 3x3x16

    reconstructed_input = self.unpool(encoded_input,indices3,output_size=to
    reconstructed_input = self.decoder_conv1(reconstructed_input) #7x7x16 to 7x7x16
    reconstructed_input = self.unpool(reconstructed_input,indices2) #7x7x16 to 14x14x8
    reconstructed_input = self.decoder_conv2(reconstructed_input)#14x14x8 to 14x14x16
    reconstructed_input = self.unpool(reconstructed_input,indices1)#14x14x16 to 28x28x8
    reconstructed_input = self.decoder_conv3(reconstructed_input)#28x28x8 to 28x28x1

    return reconstructed_input,encoded_input

```

```
In [ ]: model_5_a = AE5_ConvAE_with_unpooling()
criterion_5_a = nn.MSELoss()
optimizer_5_a = torch.optim.Adam(model_5_a.parameters(),lr=0.001)

training_loss_5_a=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        outputs,_=model_5_a(images)
        loss=criterion_5_a(outputs,images)
        training_loss_5_a.append(loss.item())

        optimizer_5_a.zero_grad()
        loss.backward()
        optimizer_5_a.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

```
In [ ]: class AE5_ConvAE_with_deconv(nn.Module):
    def __init__(self):
        super(AE5_ConvAE_with_deconv,self).__init__()

        #encoder
        self.encoder_conv1 = nn.Sequential(
            nn.Conv2d(1,8, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2))
        )
        self.encoder_conv2 = nn.Sequential(
            nn.Conv2d(8,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2))
        )
        self.encoder_conv3 = nn.Sequential(
            nn.Conv2d(16,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2))
        )

        #decoder module
        self.decoder_conv1 = nn.Sequential(
            nn.ConvTranspose2d(16,16, kernel_size = 3, stride = 2),
            nn.ReLU()
        )
        self.decoder_conv2 = nn.Sequential(
            nn.ConvTranspose2d(16,8, kernel_size = 4, stride = 2, padding = 1),
            nn.ReLU()
        )
        self.decoder_conv3 = nn.Sequential(
            nn.ConvTranspose2d(8,1, kernel_size = 4, stride = 2, padding = 1),
            nn.ReLU()
        )

    def forward(self,x):

        encoded_input = self.encoder_conv1(x.float())
        encoded_input = self.encoder_conv2(encoded_input)
        encoded_input = self.encoder_conv3(encoded_input)

        reconstructed_input = self.decoder_conv1(encoded_input)
        reconstructed_input = self.decoder_conv2(reconstructed_input)
        reconstructed_input = self.decoder_conv3(reconstructed_input)

        return reconstructed_input,encoded_input
```



```
In [ ]: model_5_b = AE5_ConvAE_with_deconv()
criterion_5_b = nn.MSELoss()
optimizer_5_b = torch.optim.Adam(model_5_b.parameters(),lr=0.001)

training_loss_5_b=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        outputs,_=model_5_b(images)
        loss=criterion_5_b(outputs,images)
        training_loss_5_b.append(loss.item())

        optimizer_5_b.zero_grad()
        loss.backward()
        optimizer_5_b.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

```

In [ ]: class AE5_ConvAE_with_deconv_unpool(nn.Module):
    def __init__(self):
        super(AE5_ConvAE_with_deconv_unpool,self).__init__()

        #encoder
        self.encoder_conv1 = nn.Sequential(
            nn.Conv2d(1,8, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        )
        self.encoder_conv2 = nn.Sequential(
            nn.Conv2d(8,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        )
        self.encoder_conv3 = nn.Sequential(
            nn.Conv2d(16,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        )

        #initializing the decoder module
        self.decoder_conv1 = nn.Sequential(
            nn.ConvTranspose2d(16,16, kernel_size = 3, stride = 1, padding = 1),
            nn.ReLU()
        )
        self.decoder_conv2 = nn.Sequential(
            nn.ConvTranspose2d(16,8, kernel_size = 3, stride = 1, padding = 1),
            nn.ReLU()
        )
        self.decoder_conv3 = nn.Sequential(
            nn.ConvTranspose2d(8,1, kernel_size = 3, stride = 1, padding = 1),
            nn.ReLU()
        )

        #unpooling
        self.unpool = nn.MaxUnpool2d(kernel_size = (2,2))

    def forward(self,x): #defines the forward pass and also the structure of the ne

        encoded_input,indices1 = self.encoder_conv1(x.float())
        encoded_input,indices2 = self.encoder_conv2(encoded_input)
        encoded_input,indices3 = self.encoder_conv3(encoded_input)

        reconstructed_input = self.unpool(encoded_input,indices3,output_size=torch.
        reconstructed_input = self.decoder_conv1(reconstructed_input)
        reconstructed_input = self.unpool(reconstructed_input,indices2)
        reconstructed_input = self.decoder_conv2(reconstructed_input)
        reconstructed_input = self.unpool(reconstructed_input,indices1)
        reconstructed_input = self.decoder_conv3(reconstructed_input)

        return reconstructed_input,encoded_input

```

```

In [ ]: model_5_c = AE5_ConvAE_with_deconv_unpool()
        criterion_5_c = nn.MSELoss()
        optimizer_5_c = torch.optim.Adam(model_5_c.parameters(),lr=0.001)

        training_loss_5_c=[]

        for epoch in range(epochs):
            for images,labels in train_loader:
                outputs,_=model_5_c(images)
                loss=criterion_5_c(outputs,images)
                training_loss_5_c.append(loss.item())

                optimizer_5_c.zero_grad()
                loss.backward()
                optimizer_5_c.step()

        print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")

```

```

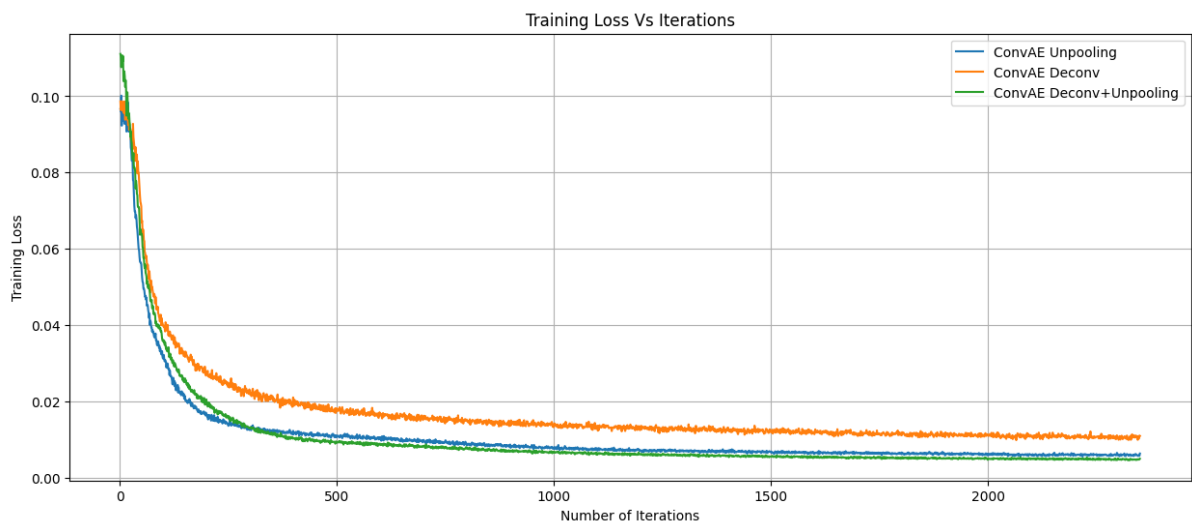
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed

```

```

In [ ]: plt.plot(range(1,len(training_loss_5_a)+1),training_loss_5_a,label="ConvAE Unpoolin
plt.plot(range(1,len(training_loss_5_a)+1),training_loss_5_b,label="ConvAE Deconv")
plt.plot(range(1,len(training_loss_5_a)+1),training_loss_5_c,label="ConvAE Deconv+U
plt.legend()
plt.grid()
plt.title("Training Loss Vs Iterations")
plt.xlabel("Number of Iterations")
plt.ylabel("Training Loss")
plt.show()

```



```

In [ ]: model_5_a.eval()
        with torch.no_grad():
            for images in test_sample_loader:
                images=images.reshape(10,1,28,28)
                outputs_5_a,_ = model_5_a(images.float())

model_5_b.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images=images.reshape(10,1,28,28)
        outputs_5_b,_ = model_5_b(images.float())

activation_5_c=[]
model_5_c.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images=images.reshape(10,1,28,28)
        outputs_5_c,_ = model_5_c(images.float())

plt.rcParams["figure.figsize"] = (15,6)

for i in range(10):
    fig, (ax1, ax2,ax3,ax4) = plt.subplots(1,4)

    ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_5_a[i].detach().numpy().reshape(28,28),cmap='gray')
    ax2.set_title('Reconstruction: CAE_Unpooling')
    ax2.axis("off")


    ax3.imshow(outputs_5_b[i].detach().numpy().reshape(28,28),cmap='gray')
    ax3.set_title('Reconstruction: CAE_Deconv')
    ax3.axis("off")

    ax4.imshow(outputs_5_c[i].detach().numpy().reshape(28,28),cmap='gray')
    ax4.set_title('Reconstruction: CAE_Deconv+Unpooling')
    ax4.axis("off")

    plt.show()

    print("Reconstruction Error in ConvAE Unpooling:",np.sum(np.dot(((images[i].detach()
    print("Reconstruction Error in ConvAE Deconv:",np.sum(np.dot(((images[i].detach()
    print("Reconstruction Error in ConvAE Deconv+Unpooling:",np.sum(np.dot(((images[i]
    print("_____


```



Reconstruction Error in ConvAE Unpooling: 11.234124905335815

Reconstruction Error in ConvAE Deconv: -1.4055003756956699


Reconstruction Error in ConvAE Deconv+Unpooling: -0.3102661686320747



Reconstruction Error in ConvAE Unpooling: 286.5323405489686

Reconstruction Error in ConvAE Deconv: 60.04885587398819


Reconstruction Error in ConvAE Deconv+Unpooling: 85.61746780524234



Reconstruction Error in ConvAE Unpooling: 179.24624768836247

Reconstruction Error in ConvAE Deconv: -21.199034076658013


Reconstruction Error in ConvAE Deconv+Unpooling: -16.239284685114132



Reconstruction Error in ConvAE Unpooling: 118.51167533133857

Reconstruction Error in ConvAE Deconv: -4.64380432908483


Reconstruction Error in ConvAE Deconv+Unpooling: 15.405881911030578



Reconstruction Error in ConvAE Unpooling: 48.152131855358235

Reconstruction Error in ConvAE Deconv: 4.073457362669939


Reconstruction Error in ConvAE Deconv+Unpooling: -3.5883763536256232



Reconstruction Error in ConvAE Unpooling: 355.24374216010676

Reconstruction Error in ConvAE Deconv: 6.777887878345329


Reconstruction Error in ConvAE Deconv+Unpooling: 68.2292085601378



Reconstruction Error in ConvAE Unpooling: 139.41802261812475

Reconstruction Error in ConvAE Deconv: 21.94462291790176


Reconstruction Error in ConvAE Deconv+Unpooling: 24.037594039694973



Reconstruction Error in ConvAE Unpooling: 162.28777739301745

Reconstruction Error in ConvAE Deconv: 6.38876997719619


Reconstruction Error in ConvAE Deconv+Unpooling: 38.435406576058824



Reconstruction Error in ConvAE Unpooling: 114.93306803095905

Reconstruction Error in ConvAE Deconv: 2.440420379285225

Reconstruction Error in ConvAE Deconv+Unpooling: 15.064461754653077



Reconstruction Error in ConvAE Unpooling: 261.43499223541403

Reconstruction Error in ConvAE Deconv: 17.66976994109564

Reconstruction Error in ConvAE Deconv+Unpooling: 52.62210097727178

```
In [ ]: model_5_a.encoder_conv1[0].weight.detach().numpy().squeeze().shape
```

```
Out[ ]: (8, 3, 3)
```

```
In [ ]: #Function for visualisation of weights
from torchvision import utils

def visTensor(tensor, ch=0, allkernels=False, nrow=8, padding=1):
    n,c,w,h = tensor.shape

    if allkernels: tensor = tensor.view(n*c, -1, w, h)
    elif c != 3: tensor = tensor[:,ch,:,:].unsqueeze(dim=1)

    rows = np.min((tensor.shape[0] // nrow + 1, 64))
    grid = utils.make_grid(tensor, nrow=nrow, normalize=True, padding=padding)
    plt.figure( figsize=(nrow,rows) )
    plt.imshow(grid.numpy().transpose((1, 2, 0)))
```

```
In [ ]: #VISUALISING DECODER WEIGHTS FOR CONVOLUTION AUTOENCODER WITH UNPOOLING
```

```
filter = model_5_a.decoder_conv2[0].weight.data.clone()  
visTensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')  
plt.ioff()  
plt.title('decoder_conv2 Weights')  
plt.show()
```

```
filter = model_5_a.decoder_conv3[0].weight.data.clone()  
visTensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')  
plt.ioff()  
plt.title('decoder_conv3 Weights')  
plt.show()
```

decoder_conv2 Weights



decoder_conv3 Weights




```
In [ ]: #VISUALISING DECODER WEIGHTS FOR CONVOLUTION AUTOENCODER WITH DECONVOLUTION
```

```
filter = model_5_b.decoder_conv1[0].weight.data.clone()  
visTensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')  
plt.ioff()  
plt.title('decoder_conv1 Weights')  
plt.show()
```


```
filter = model_5_b.decoder_conv2[0].weight.data.clone()  
visTensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')  
plt.ioff()  
plt.title('decoder_conv2 Weights')  
plt.show()
```


```
filter = model_5_b.decoder_conv3[0].weight.data.clone()  
visTensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')  
plt.ioff()  
plt.title('decoder_conv3 Weights')  
plt.show()
```


decoder_conv1 Weights



decoder_conv2 Weights



decoder_conv3 Weights



```
In [ ]: #VISUALISING DECODER WEIGHTS FOR CONVOLUTION AUTOENCODER WITH DECONVOLUTION+UNPOOLING

filter = model_5_c.decoder_conv1[0].weight.data.clone()
visTensor(filter, ch=0, allkernels=False)

plt.axis('off')
plt.ioff()
plt.title('decoder_conv1 Weights')
plt.show()


filter = model_5_c.decoder_conv2[0].weight.data.clone()
visTensor(filter, ch=0, allkernels=False)

plt.axis('off')
plt.ioff()
plt.title('decoder_conv2 Weights')
plt.show()


filter = model_5_c.decoder_conv3[0].weight.data.clone()
visTensor(filter, ch=0, allkernels=False)

plt.axis('off')
plt.ioff()
plt.title('decoder_conv3 Weights')
plt.show()
```


decoder_conv1 Weights



decoder_conv2 Weights



decoder_conv3 Weights



Observations:

- Decoder weights of Deconvolution with Unpooling and Unpooling are smaller than the only Deconvolution one. This can be because, Unpooling handles part of the upsampling, reducing the need for the deconvolution layers to learn complex upsampling transformations.
- By looking at reconstruction error, it looks like unpooling and deconv+ unpooling does better on reconstruction of images.
- Visually, reconstructed images using Deconvolution are appealing.